

Lecture 14

Binary Search Trees

EECS 281: Data Structures & Algorithms

Search

- Retrieval of a particular piece of information from large volumes of previously stored data
- Purpose is typically to access information within the item (not just the key)
- Recall that arrays, linked lists are worst-case $O(n)$ for either searching or inserting

Need a data structure with optimal efficiency for searching and inserting

Symbol Table

Definition: abstract data structure of items with keys that supports two basic operations:

- *Insert* a new item
- *Search* for an item with a given key

Symbol Table: ADT

- **insert** a new item
- **search** for an item (items) with a given key
- **remove** an item with a specified key
- **sort** the symbol table
- **select** the item with the k^{th} largest key
- **join** two symbol tables


Also may want construct, test if empty, destroy, copy...

Binary Search Tree

- A binary search tree is organized as a binary tree
- The keys in a binary search tree satisfy the binary-search-tree property
 - The key of any node is:
 - keys of all nodes in its left subtree and
 - keys of all nodes in its right subtree
- Essential property of BST is that **insert** is as easy to implement as **search**

Binary Search Tree Property

- The key of any node is:
keys of all nodes in its left subtree and
keys of all nodes in its right subtree



Concrete Implementation

- Node in a Binary Tree

```
template <typename KEY>
struct Node {
    KEY key;
    Node *left, *right;
};
```

- What if we need to remove a node or add a node?
- Do we need to move up the tree?
- How would you change this implementation?

Concrete Implementation


- Node in a Binary Tree (with parent)

```
template <typename KEY>
struct Node {
    KEY key;
    Node *left, *right, *parent;
};
```

- Parent pointer is helpful for operations on internal nodes, for example, delete

Exercise

- Write the output for inorder, preorder and post-order traversals of these BSTs



```
void inorder(Node *x) {  
    if (!x) return;  
    inorder(x->left);  
    print(x->key);  
    inorder(x->right);  
} // inorder()
```

```
void preorder(Node *x) {  
    if (!x) return;  
    print(x->key);  
    preorder(x->left);  
    preorder(x->right);  
} // preorder()
```

```
void postorder(Node *x) {  
    if (!x) return;  
    postorder(x->left);  
    postorder(x->right);  
    print(x->key);  
} // postorder()
```

Sort: Binary Search Tree

Can you think of an easy method of sorting using a binary search tree?

Search

- How can we find a key in a binary search tree?

```
//return a pointer to a node with key k if  
//one exists; otherwise, return nullptr  
Node *tree_search(Node *x, Key k);
```

- Hint: remember the key of any node is:
 keys of all nodes in its left subtree and
 keys of all nodes in its right subtree
- Bonus: what are the average and worst-case complexities?

Search

```
1 //return a pointer to a node with key k if
2 //one exists; otherwise, return nullptr
3 Node *tree_search(Node *x, Key k) {
4     while (x != nullptr && k != x->key) {
5         if (k < x->key)
6             x = x->left;
7         else
8             x = x->right;
9     } // while
10    return x;
11 } // tree_search()
```


Search

```
1 //return a pointer to a node with key k if
2 //one exists; otherwise, return nullptr
3 Node *tree_search(Node *x, Key k) {
4     if (x == nullptr || x->key == k) return x;
5     if (k < x->key)
6         return tree_search(x->left, k);
7     return tree_search(x->right, k);
8 } // tree_search()
```

- Same as BST
- To search for a key k , we trace a downward path starting at the root
- The next node visited depends on the outcome of the comparison of k with the key of the current node

Search Example

```
tree_search(<ptr to 15>, 9);
```



Search


- Complexity is $O(h)$, where h is the (maximum) height of the tree
- Average case complexity: $O(\log n)$
 - Balanced tree
- Worst case complexity: $O(n)$
 - "Stick" tree

Insert

- How do we insert a new key into the tree?
- Similar to search
- Start at the root, and trace a path downwards, looking for a null pointer to append the node


Insert Example

```
tree_insert(10);
```



Insert with Duplicates

- For sets with no duplicates, use ($<$, $>$)
- For duplicates, need deterministic policy
 - Use (\leq , $>$) or ($<$, \geq)



uses ($<$, \geq)

Insert

```
1 void tree_insert(Node *&x, Key k) {  
2     if (x == nullptr) {  
3         x = new Node;  
4         x->key = k;  
5         x->left = x->right = nullptr;  
6     } // if  
7     else if (k < x->key)  
8         tree_insert(x->left, k);  
9     else  
10        tree_insert(x->right, k);  
11 } // tree_insert()
```

- New node inserted at leaf
- Note the nifty(?) use of reference-to-pointer-to-Node
- Exercise: modify this code to set the parent pointer

Insert

```
1 void tree_insert(Node *&x, Node *x_parent, Key k) {  
2     if (x == nullptr) {  
3         x = new Node;  
4         x->key = k;  
5         x->left = x->right = nullptr;  
6         x->parent = x_parent;  
7     } // if  
8     else if (k < x->key)  
9         tree_insert(x->left, x, k);  
10    else  
11        tree_insert(x->right, x, k);  
12 } // tree_insert()
```

- This version sets the parent pointer

Insert [CLRS]

```
1 void tree_insert_CLRS(Node *&t, Key k) {
2     Node *y = nullptr;
3     Node *x = t;
4     while (x != nullptr) { //find location for new node
5         y = x;
6         x = (k < x->key) ? x->left : x->right;
7     } // while
8     Node *z = new Node(k, y); //y is parent of new node
9
10    if (y == nullptr)
11        t = z; //tree t was empty
12    else if (z->key < y->key) y->left = z;
13    else y->right = z;
14 }
```

Exercise

- Start with an empty tree
- Insert these keys, in this order:
12, 5, 18, 2, 9, 15, 19, 17, 13
- Draw the tree
- Write a new order to insert the same keys which generates a worst-case tree
- How many worst-case case trees are possible for n nodes?


Complexity

- The complexity of insert (and many other tree functions) depends on the height of the tree
- Average case (balanced): $O(\log n)$
- Worst case (unbalanced "stick"): $O(n)$
- Average case:
 - Random data
 - Likely to be well-balanced

Exercise

- Write a function to find the Node with the smallest key
- What are the average and worst-case complexities?

//returns a pointer to the Node with the min key
`Node *tree_min(Node *x);`



Exercise

```
1 //return a pointer to the min key node
2 Node *tree_min(Node *x) {
3     if (x == nullptr)
4         return nullptr;
5     while (x->left)
6         x = x->left;
7     return x;
8 } // tree_min()
```


- Average case complexity: $O(\log n)$
- Worst case complexity: $O(n)$

Delete

- What if we want to delete a node?
- To delete node z :
 1. z has no children (trivial)
 2. z has no left child
 3. z has no right child
 4. z has two children
- Complete algorithm is in CLRS 12.3


Delete

- z has no left child: replace z by right child



Delete


- z has no right child: replace z by left child



Delete


- z has left and right children
- Replace with a “combined” tree of both
- Key observation
 - All in LHS subtree all in RHS subtree
 - Transplant smallest RHS node to root
 - Called the inorder successor
 - Must be some such node, since RHS is not empty
 - New root might have a right child, but no left child
 - Make new root’s left child the LHS subtree

Delete




Delete

1. Transplant smallest RHS node to root
 - Must be some such node, since RHS is not empty
 - New root might have a right child, but no left child
2. Make new root's left child the LHS subtree




Delete

- This is where parent pointers are especially helpful: the "transplant" is $O(1)$
- What is the complexity of finding the smallest node in the right subtree?




Delete

- This is where parent pointers are especially helpful: the "transplant" is $O(1)$
- What is the complexity of finding the smallest node in the right subtree?
 - $O(n)$ worst case, $O(\log n)$ average case, which dominates the complexity of delete



Delete

- Easier case: z 's right child has no left child



Delete Helper Function [CLRS]

```
1 void tree_transplant(Node *&t, Node *u, Node *v) {  
2     if (u->parent == nullptr)  
3         t = v;  
4     else if (u == u->parent->left)  
5         u->parent->left = v;  
6     else  
7         u->parent->right = v;  
8  
9     if (v != nullptr)  
10        v->parent = u->parent;  
11    delete u;  
12 } // tree_transplant()
```

Delete [CLRS]

```
1 void tree_delete(Node *&t, Node *z) {
2     if (z->left == nullptr)
3         tree_transplant(t, z, z->right);
4     else if (z->right == nullptr)
5         tree_transplant(t, z, z->left);
6     else {
7         Node *y = tree_min(z->right);
8         if (y->parent != z) {
9             tree_transplant(t, y, y->right); // Don't delete!
10            y->right = z->right;
11            y->right->parent = y;
12        } // if
13        tree_transplant(t, z, y);
14        y->left = z->left;
15        y->left->parent = y;
16    } // else
17 } // tree_delete()
```

Single-Function Delete 1/3

```
1  template <typename T>
2  void BinaryTree<T>::remove(Node *&tree, const T &val)
3  {
4      Node *nodeToDelete = tree;
5      Node *inorderSuccessor;
6
7      // Recursively find the node containing the value to delete
8      if (tree == nullptr)
9          return;
10     else if (val < tree->value)
11         remove(tree->left, val);
12     else if (tree->value < val)
13         remove(tree->right, val);
14     else {
```

Single-Function Delete 2/3

```
15      // Check for simple cases where one subtree is empty
16      if (tree->right == nullptr)
17      {
18          tree = tree->left;
19          delete nodeToDelete;
20      } // if
21      else if (tree->left == nullptr)
22      {
23          tree = tree->right;
24          delete nodeToDelete;
25      } // else if
```

Single-Function Delete 3/3

```
26         else {
27             // Node to delete has both left and right subtrees
28             inorderSuccessor = tree->right;
29
30             while (inorderSuccessor->left != nullptr)
31                 inorderSuccessor = inorderSuccessor->left;
32
33             // Replace value with one from inorder successor
34             nodeToDelete->value = inorderSuccessor->value;
35             // Remove the inorder successor from right subtree
36             remove(tree->right, inorderSuccessor->value);
37         } // else
38     } // else
39 } // BinaryTree::remove()
```

Summary: Binary Search Trees

- Each node points to two children (left, right), and possibly a parent
- All nodes are ordered: left root right
- Modification of nodes
 - External is easy
 - Internal is more complicated
- In general, operations on BSTs are:
 - $O(\log n)$ average
 - $O(n)$ worst case