

Lecture 16

Hash Collision Resolution

EECS 281: Data Structures & Algorithms

Collision Resolution

Definition: method to handle case when two keys hash to same address

Methods of Collision Resolution

- Separate Chaining
- Linear Probing
- Quadratic Probing
- Double Hashing

Collision Resolution

Separate Chaining: scheme for collision resolution where we maintain M linked lists, one for each table address

Collision Resolution

Property: Separate chaining reduces the number of comparisons for sequential search by a factor of M (on average), using extra space for M links

Property: In a separate chaining hash table with M lists (table addresses) and N keys, the probability that the number of keys in each list is within a small constant factor of N/M is extremely close to 1 ($O(1)$) *if the hash function is good*

Collision Resolution

- Separate chaining
 - Insert: constant time
 - $O(1)$
 - Search: time proportional to N/M
 - $O(N/M)$
 - Remove: dependent upon Search
 - $O(N/M)$

This is why we choose $M \approx N$: $O(N/M) = O(1)$

Collision Resolution

Use empty places in table to resolve collisions (known as *open-addressing*)

Probe: determination whether given table location is 'occupied'

Linear Probing: when collision occurs, check the next position in the table

Possible Probe Outcomes

- Miss: probe finds empty cell in table, OR
- Hit: probe finds cell that contains item whose key matches search key, OR
- Full: probe finds cell has 'occupant', but key doesn't match search key
- *If probe results in full, then probe table at next "higher" cell until hit (search ends successfully) or miss (search ends unsuccessfully)*

Cluster

Definition: contiguous group of occupied table cells

Consider table that is half-full ($M = 2N$)

What is best case/worst case distribution?

- Best Case:
- Worst Case:

Cluster

Consider table that is half-full ($M = 2N$)

Pop Quiz

- What is the *average* cost (in terms of N) to obtain a miss (find an empty cell) given the best case distribution?
- What is the *average* cost (in terms of N) to obtain a miss (find an empty cell) given the worst case distribution?

Linear Probing

- How to delete a key from a table built with linear probing?
 - Why is this hard?
- Option 1: remove it, re-hash rest of cluster
- Option 2: use a “dummy” element
 - Not an element, not empty either
 - We’ll call this ‘deleted’

“Deleted” Elements

- When an item is marked as deleted:
 - Insert considers it an empty spot and usable
 - Search considers it occupied (full) and keeps searching

Possible Probe Outcomes (Revised)

- Empty: probe finds cell that has never held item, OR
- Deleted: probe finds cell that once held item, but is not currently holding item, OR
- Hit: probe finds cell that contains item whose key matches search key, OR
- Full: probe finds cell has 'occupant', but key doesn't match search key

Load Factor (α)

- $\alpha = N/M$, where N keys are placed in an M -sized table
- Separate Chaining
 - α is average number of items per list
 - α is sometimes larger than 1
- Linear Probing
 - α is percentage of table positions occupied
 - α is (must be) ≤ 1

Collision Resolution

When collisions are resolved with linear probing, the average number of probes required to search in a hash table of size M that contains $N = \alpha M$ keys is about

$$\frac{1}{2} \left(1 + \frac{1}{1 - \alpha} \right) \div \quad \text{for hits}$$

$$\frac{1}{2} \left(1 + \frac{1}{(1 - \alpha)^2} \right) \div \quad \text{for misses}$$

Effect of Load on # of Probes

Linear Probing

α (% Full)	Average Probes: Successful Search	Average Probes: Unsuccessful Search
.1	1.1	1.1
.2	1.1	1.3
.3	1.2	1.5
.4	1.3	1.9
.5	1.5	2.5
.6	2.8	3.6
.7	2.2	6.1
.8	3.0	13.0
.9	5.5	50.5

Collision Resolution

Quadratic Probing

Try buckets at increasing 'distance' from hash table location

- $h(key) \bmod M \Rightarrow addr$
- if bucket $addr$ is full, then try
 - $(h(key) + j^2) \bmod M$ for $j = 1, 2, \dots$

Collision Resolution

When collisions are resolved with quadratic probing, the average number of probes required to search in a hash table of size M that contains $N = \alpha M$ keys is about

$$\frac{1}{2} + \ln \frac{1}{1-\alpha} \quad \div \quad \text{for hits}$$

$$\frac{1}{1-\alpha} + \ln \frac{1}{1-\alpha} \quad \div \quad \text{for misses}$$

Effect of Load on # of Probes

Quadratic Probing

α (% Full)	Average Probes: Successful Search	Average Probes: Unsuccessful Search
.1	1.06	1.12
.2	1.12	1.27
.3	1.21	1.49
.4	1.31	1.78
.5	1.44	2.19
.6	1.62	2.82
.7	1.85	3.84
.8	2.21	5.81
.9	2.85	11.40

Collision Resolution

Double Hashing

Apply additional hash function if collision occurs

- $h(\text{key}) \bmod M \Rightarrow \text{addr}$
- If bucket *addr* is full, then try
 - $(h(\text{key}) + j * h'(\text{key})) \bmod M$, where
 - $j = 1, 2, 3, \dots$ and
 - $h'(k) = q - (k \bmod q)$ for some prime number $q < M$
 - Until an empty cell is found

New Topic: Dynamic Hashing

- As number of keys in hash table increases, search performance degrades
- Separate Chaining
 - Search time increases gradually
 - Double keys means double list length at each of M table locations
- Linear Probing
 - Search time increases dramatically as table fills
 - May reach point when no more keys can be inserted

Objective: Dynamic Hashing

- Double size of table when it 'fills up' (load factor is say one half)
- Expensive, but infrequent

Amortized Analysis

- Cannot guarantee that each and every operation will be fast, but can guarantee that average cost per operation will be low
- Total cost is low, but performance profile is erratic
- Most operations are extremely fast, but some operations require much more time to rebuild the table

Amortized Analysis: Concept

- Each insert
 - Pays (small constant) cost to actually insert
 - Deposits other small constant (“balance”) in a bank
- First $M/2 - 1$: build up “balance”
- $(M/2)$ th insertion
 - Faced with a big (not small constant) bill
 - Finds a big (not small constant) balance
- Net result
 - Each insert charged small constant costs
 - Some costs deferred

Amortized Analysis: Applied

- Start with table of size M
- Each insertion in a table $\leq \frac{1}{2}$ full
 - Costs up to 2.5 probes (from table)
- Insert $M/2 - 1$ keys
 - $2.5 * (M/2 - 1) = O(M)$

Amortized Analysis: Applied

- Insert $(M/2)^{\text{th}}$ key
- Build new table, size $2M$
 - Remove keys from old table, insert in new
 - Each insert $\leq \frac{1}{4}$ full, costs a maximum of 1.5 probes (from table)
 - $1.5 * M/2 = O(M)$
- $O(M) + O(M) = O(M)$
 - Linear time to insert $M/2$ keys, but last one is at a higher cost than the previous inserts

Summary: Hashing

- Collision Resolution
 - Separate Chaining creates a linked list for each table address
 - Linear Probing uses empty places in table to resolve collisions
 - Quadratic Probing looks for empty table address at increasing distance from original hash
 - Double Hashing applies additional hash function to original hash
- Dynamic Hashing increases the table size when it reaches some pre-determined load factor